

ON CONSTRUCTING MINIMUM SPANNING TREES IN k -DIMENSIONAL SPACES AND RELATED PROBLEMS*

ANDREW CHI-CHIH YAO†

Abstract. The problem of finding a minimum spanning tree connecting n points in a k -dimensional space is discussed under three common distance metrics: Euclidean, rectilinear, and L_∞ . By employing a subroutine that solves the post office problem, we show that, for fixed $k \geq 3$, such a minimum spanning tree can be found in time $O(n^{2-a(k)}(\log n)^{1-a(k)})$, where $a(k) = 2^{-(k+1)}$. The bound can be improved to $O((n \log n)^{1.8})$ for points in 3-dimensional Euclidean space. We also obtain $o(n^2)$ algorithms for finding a farthest pair in a set of n points and for other related problems.

Key words. algorithm, minimum spanning tree, nearest neighbor, post office problem

1. Introduction. Given an undirected graph with a weight assigned to each edge, a minimum spanning tree (MST) is a spanning tree whose edges have a minimum total weight among all spanning trees. The classical algorithms for finding MST were given by Dijkstra [7], Kruskal [13], Prim [14], and Sollin [4, p. 179]. It is well known (e.g., see Aho, Hopcroft and Ullman [1]) that, for a graph with n vertices, an MST can be found in $O(n^2)$ time. (All time bounds discussed in this paper are for the worst-case behavior of algorithms.) For a sparse graph with e edges and n vertices, it was shown by Yao [16] that an MST can be found in time $O(e \log \log n)$. More studies of MST algorithms can also be found in Cheriton and Tarjan [6], Kerschenbaum and Van Slyke [11].

An interesting application of MST occurs in connection with hierarchical clustering analysis in pattern recognition (see, for example, Dude and Hart [9, Ch. 6], Zahn [20]). In this application, n vertices $V = \{\tilde{v}_1, \tilde{v}_2, \dots, \tilde{v}_n\}$ are given, each a k -tuple of numbers. The graph is understood to be a complete graph $G(V)$ on these n vertices, with the weight on each edge $\{\tilde{v}_i, \tilde{v}_j\}$ being $d(\tilde{v}_i, \tilde{v}_j)$ where d is a certain metric function computable from the components of \tilde{v}_i and \tilde{v}_j . A simple way to find an MST in this case is to compute all the weights $d(\tilde{v}_i, \tilde{v}_j)$, and then use an $O(n^2)$ MST algorithm for general graphs. However, as there are only kn input parameters, it is interesting to find out if there are algorithms which take only $o(n^2)$ time. Several empirically good algorithms were proposed in Bentley and Friedman [2], where a list of references to other applications of finding MST in k -dimensional spaces can also be found. Shamos and Hoey [16] gave an $O(n \log n)$ algorithm for n points in the plane ($k = 2$) with Euclidean metric. No algorithm, however, is known to have a guaranteed bound of $o(n^2)$ when $k \geq 3$.

In this paper, we consider three common metrics in k -dimensional spaces, namely, the rectilinear (L_1), the Euclidean (L_2), and the L_∞ metric. We use E_p^k ($p = 1, 2, \infty$) to denote the space of all k -tuples of real numbers with the L_p -metric, i.e., the distance between two points \tilde{x} and \tilde{y} is given by $d_p(\tilde{x}, \tilde{y}) = (\sum_{i=1}^k |x_i - y_i|^p)^{1/p}$. (It is agreed that $d_\infty(\tilde{x}, \tilde{y}) = \max_i |x_i - y_i|$.) We give new algorithms which construct, for a given set V of n points in E_p^k , an MST for the associated complete graph $G(V)$. The algorithms work in time $O(n^{2-a(k)}(\log n)^{1-a(k)})$, where $a(k) = 2^{-(k+1)}$ for any fixed $k \geq 3$. Fast algorithms for related geometric problems are also given using similar techniques.

* Received by the editors November 3, 1980. This research was supported in part by National Science Foundation under grant MCS 72-03752 A03.

† Computer Science Department, Stanford University, Stanford, California 94305.

The main results of this paper are summarized in the following theorem. Sections 2–5 are devoted to a proof of it.

THEOREM 1. *Let $k \geq 3$ be a fixed integer, $a(k) = 2^{-(k+1)}$, and all points to be considered are in E_p^k with $p \in \{1, 2, \infty\}$. Then each of the following problems can be solved in time $O(n^{2-a(k)}(\log n)^{1-a(k)})$. For the case when $k = 3$ and $p = 2$, the bound can be improved to $O((n \log n)^{1.8})$.*

MST-problem	Let V be a set of n points; find a minimum spanning tree on V .
NFN-problem	(nearest foreign neighbor). Let V_1, V_2, \dots, V_l be disjoint sets of points, $V = \bigcup_i V_i$, and $ V = n$. For each V_i and every $\tilde{x} \in V_i$, find a $\tilde{y} \in V - V_i$ such that $d_p(\tilde{x}, \tilde{y}) = \min \{d_p(\tilde{x}, \tilde{z}) \mid \tilde{z} \in V - V_i\}$.
GN-problem	(geographic neighbor). Let V be a set of n points. For any $\tilde{x} \in V$, let $N(\tilde{x}) = \{\tilde{v} \mid v_i \cong x_i \text{ for all } 1 \leq i \leq k, \tilde{v} \neq \tilde{x}, \tilde{v} \in V\}$. For each $\tilde{x} \in V$, find a $\tilde{y} \in N(\tilde{x})$ such that $d_p(\tilde{x}, \tilde{y}) = \min \{d_p(\tilde{x}, \tilde{v}) \mid \tilde{v} \in N(\tilde{x})\}$ if $N(\tilde{x}) \neq \emptyset$.
AFP-problem	(all farthest points) [3]. Let V be a set of n points. For each $\tilde{x} \in V$, find a $\tilde{y} \in V$ such that $d_p(\tilde{x}, \tilde{y}) = \max \{d_p(\tilde{x}, \tilde{v}) \mid \tilde{v} \in V\}$.
FP-problem	(farthest pair) [3]. Let V be a set of n points; find $\tilde{x}, \tilde{y} \in V$ such that $d_p(\tilde{x}, \tilde{y}) = \max \{d_p(\tilde{u}, \tilde{v}) \mid \tilde{u}, \tilde{v} \in V\}$.

In § 6, we briefly describe, for the L_2 and the L_∞ metric, how to obtain $o(kn^2)$ algorithms when k is allowed to vary with n .

A remark on the model of computation: We assume a random access machine with arithmetic on real numbers, and charge uniform cost for all access and arithmetic operations [1]. In this paper, we often carry out computations of $d_p(\tilde{x}, \tilde{y})$, which involves an apparent square root operation when $p = 2$. However, since our construction of MST depends only on the linear ordering among the edge weights, we can replace $d_p(\tilde{x}, \tilde{y})$ throughout by some monotone function of $d_p(\tilde{x}, \tilde{y})$. In particular, $d_2(\tilde{x}, \tilde{y})$ may be replaced by $(d_2(\tilde{x}, \tilde{y}))^2 = \sum (x_i - y_i)^2$ everywhere to produce a valid algorithm without square root operations. We shall, however, retain the original form of the algorithm for clarity and for consistency with the cases $p = 1, \infty$.

2. The post-office problem and its applications. In this section we review solutions to the post-office problem, and show how it can be used to prove Theorem 1 for the AFP, FP and NFN problems.

The *post-office problem* can be stated as follows. Given a set of n points $\tilde{v}_1, \tilde{v}_2, \dots, \tilde{v}_n$ in E_p^k , we wish to preprocess them so that any subsequent query of the following form can be answered quickly:

nearest-point query. Given a point \tilde{x} , find a nearest \tilde{v}_i to \tilde{x} (i.e., $d_p(\tilde{x}, \tilde{v}_i) \leq d_p(\tilde{x}, \tilde{v}_j)$ for all j).

This problem was mentioned in Knuth [12] for the case of points in the Euclidean plane ($k = p = 2$). For this special case, several solutions were given by Dobkin and Lipton [8] and Shamos [15]. For example, it is known that with an $O(n^2)$ -time preprocessing, any nearest-point query can be answered in $O(\log n)$ time [15]. A solution for the k -dimensional Euclidean space was given in Dobkin and Lipton [8], where it was shown that it is possible to preprocess n points such that any subsequent nearest-point query can be answered in $O(2^k \log n)$ time. Their technique is quite general, and applies equally well if we wish to answer “farthest-point” queries—Given \tilde{x} , find a farthest \tilde{v}_i to \tilde{x} —instead of nearest-point queries. The preprocessing procedure was not discussed in great detail in [8]. A detailed study [19] gives the following result.

DEFINITION. We shall use $b(k) = 2^{k+1}$, and $a(k) = b(k)^{-1} = 2^{-(k+1)}$.

LEMMA 2.1. *Let $k \geq 3$ be a fixed integer, and $p \in \{1, 2, \infty\}$. There is an algorithm which preprocesses n points in E_p^k in time $O(n^{b(k)})$ such that each subsequent nearest-point query can be answered in $O(\log n)$ time. In the special case $k = 3, p = 2$, the preprocessing time can be improved to $O(n^5 \log n)$ with a query response-time $O((\log n)^2)$. The preceding statements remain true if the farthest-point query is used in place of the nearest-point query.*

We shall now demonstrate the use of Lemma 2.1 by applying it to solve the MST problem in a special case. It also gives us some insight into the connection between MST and some typical nearest-neighbor problem [3], [16].

Consider the case when V consists of two widely separated clusters A and B . For definiteness, assume that $d_p(A, B) > n \cdot (\text{diam}(A) + \text{diam}(B))^1$. In this case any MST on V consists of the union of an MST for A and an MST for B , plus a shortest edge between A and B . Thus, to be able to solve the MST problem efficiently, we have to be able to solve the following problem efficiently:

Problem RMST. Given two well-separated sets A and B in E_p^k , with $|A| = |B| = n$, find a shortest edge between A and B .

This problem looks very similar to the problem of finding the closest pair in a set, which has an $O(n \log n)$ -time algorithm. However, there does not seem to be any simple divide-and-conquer $o(n^2)$ solution. We shall presently give an $o(n^2)$ -time algorithm employing the post-office problem as a subroutine.

Consider the following algorithm.

- (S1) Divide B into $r = \lceil n/q \rceil$ sets B_1, B_2, \dots, B_r , each with at most q points (q to be determined).
- (S2) For each $1 \leq i \leq r$, preprocess B_i for nearest-point queries as in Lemma 2.1.
- (S3) For each $\tilde{x} \in A$ and each $1 \leq i \leq r$, find a point $\tilde{y}(\tilde{x}, i) \in B_i$ that is nearest to \tilde{x} among all points in B_i .
- (S4) For each $\tilde{x} \in A$, find a $\tilde{z}(\tilde{x}) \in B$ nearest to x by comparing $\tilde{y}(\tilde{x}, i)$ for all $1 \leq i \leq r$.
- (S5) Find a shortest such edge $\{\tilde{x}, \tilde{z}(\tilde{x})\}$.

The time taken is dominated by (S2) and (S3), i.e.,

$$O(r \cdot q^{b(k)} + nr \log q).$$

Choosing $q = (\log n)^{(b(k))^{-1}}$ gives the time $O(n(n \log n)^{1-(b(k))^{-1}})$. Thus, we have found an algorithm that solves RMST in time $O(n^{2-a(k)}(\log n)^{1-a(k)})$. For the case $k = 3$ and $p = 2$, one can choose $q = (n \log n)^{1/5}$ to obtain an $O((n \log n)^{1.8})$ algorithm.

We wish to make two observations concerning the above procedure. Firstly, the AFP- and FP-problems can be solved with the same time bounds by very similar procedures (employing farthest-point queries and preprocessing, of course). We will thus consider that Theorem 1 has been proved for these problems. Secondly, the RMST-problem is a type of nearest-neighbor problem with some restrictions on the "legal" neighbors. It is reasonable to expect more such problems can be solved with similar techniques. The NFN- and GN-problems are problems of this type, and we will see that their efficient solutions enable the MST problem to be solved efficiently. We shall give a fast algorithm for NFN-problems presently, leaving the more involved proof of Theorem 1 for MST and GN to the later sections.

¹ We use the notation $d_p(A, B) = \min \{d_p(\tilde{u}, \tilde{v}) \mid \tilde{u} \in A, \tilde{v} \in B\}$, $d_p(\tilde{u}, S) = \min \{d_p(\tilde{u}, \tilde{v}) \mid \tilde{v} \in S\}$, and $\text{diam}(S) = \max \{d_p(\tilde{u}, \tilde{v}) \mid \tilde{u}, \tilde{v} \in S\}$.

We are given disjoint sets V_1, V_2, \dots, V_l with a total of n points in $V = \bigcup_i V_i$. For a point $\tilde{x} \in V_i$, every point $\tilde{y} \in V - V_i$ is a *foreign neighbor* of \tilde{x} . Let $q = \lceil (n \log n)^{a(k)} \rceil$; call a set V_i *small* if $|V_i| < q$, and *large* if $|V_i| \geq q$. We partition V into $r = O(n/q)$ parts B_1, B_2, \dots, B_r , where each part (call it a *block*) either is the union of several small V_i or is totally contained in some large V_i . Furthermore, each part contains at most $2q$ points, and except possibly for B_r , at least q points. The above partition can be accomplished in $O(n)$ time by breaking each large V_i into several blocks and grouping small V_i into blocks of appropriate sizes. We now preprocess each block B_i so that, for any query point \tilde{x} , a point nearest to \tilde{x} in B_i can be found in $O(\log q)$ time. According to Lemma 2.1, this preprocessing can be accomplished in time $O(rq^{b(k)})$ for all blocks B_i . We are now ready to find, for each point $\tilde{x} \in V$, a nearest foreign neighbor \tilde{y} , i.e., $d_p(\tilde{x}, \tilde{y}) = \min \{d_p(\tilde{x}, \tilde{z}) \mid \tilde{z} \in V - V_i\}$, when $\tilde{x} \in V_i$. Assume that $\tilde{x} \in V_i$ and $\tilde{x} \in B_r$. Let us find, for each block B_j that is disjoint from V_i , a point $\tilde{z}(\tilde{x}, j)$ nearest to \tilde{x} among all points in B_j . Then we find a nearest foreign neighbor \tilde{y} from the points $\tilde{z}(\tilde{x}, j)$ and points in $B_i - V_i$ by computing and comparing their distances to \tilde{x} . The running time for finding \tilde{y} , for each \tilde{x} , is thus $O(r \log q + (r + q))$. In summary, the total running time of the above procedure for NFN is $O(n + rq^{b(k)} + nr \log q + nq)$, which is $O(n^{2-a(k)}(\log n)^{1-a(k)})$. As before, an $O((n \log n)^{1.8})$ algorithm can be obtained for the case $k = 3$ and $p = 2$.

This proves Theorem 1 for the NFN-problem. An interesting connection exists between MST- and NFN-problems. In fact, in Sollin's algorithm[4, p. 179], an MST can be found essentially by solving NFN-problems $O(\log n)$ times. Thus, we have shown that an MST can be found in $\log n \times O(n^{2-a(k)}(\log n)^{1-a(k)})$ -time. The $\log n$ factor can be avoided by reducing MST to a generalized version of the GN-problem, which can be solved in time $O(n^{2-a(k)}(\log n)^{1-a(k)})$. The proof requires additional techniques beyond the simple application of post-office problems to small parts of V . We shall illustrate the ideas for two dimensions in the next section, and complete the proof in §§ 4 and 5.

3. An illustration in two dimensions. We illustrate the ideas of our MST-algorithms with an informal description for 2-dimensional Euclidean space. Let us first consider a special type of nearest-neighbor problem. Let \tilde{p} be any point in the plane. We divide the plane into eight regions relative to \tilde{p} as shown in Fig. 1. The regions are formed by four lines passing through \tilde{p} and having angles of $0^\circ, 45^\circ, 90^\circ$, and 135° , respectively, with the x -axis. We number the regions counterclockwise as shown in Fig. 1, and use $R_l(\tilde{p})$ to denote the set of points in the l th region (including its boundary), for $1 \leq l \leq 8$.


LEMMA 3.1. *If \tilde{q} and \tilde{q}' are two points in $R_l(\tilde{p})$ for some l , then $d_2(\tilde{q}, \tilde{q}') < \max \{d_2(\tilde{p}, \tilde{q}), d_2(\tilde{p}, \tilde{q}')\}$.*

Proof. Consider the triangle $\tilde{p}\tilde{q}\tilde{q}'$ (see Fig. 1). Since $\angle \tilde{q}\tilde{p}\tilde{q}' \leq 45^\circ < \pi/3$, its opposite side $\tilde{q}\tilde{q}'$ cannot be the longest side of the triangle. \square

Let V be a set of n distinct points in the plane. For each point $\tilde{v} \in V$, let $N_l(\tilde{v})$ be those points of V , excluding \tilde{v} itself, that are in the l th region relative to \tilde{v} . That is,

$$N_l(\tilde{v}) = V \cap R_l(\tilde{v}) - \{\tilde{v}\} \quad \text{for } 1 \leq l \leq 8.$$

A point \tilde{u} in $N_l(\tilde{v})$ is said to be a *nearest neighbor to \tilde{v} in the l th region* if $d_2(\tilde{v}, \tilde{u}) = \min \{d_2(\tilde{v}, \tilde{w}) \mid \tilde{w} \in N_l(\tilde{v})\}$. Note that such a nearest neighbor does not exist if $N_l(\tilde{v}) = \emptyset$, and may not be unique when it exists. Now, consider the following computational problem.


 FIG. 1. Regions $R_l(\tilde{p})$ for $1 \leq l \leq 8$.

The eight-neighbors problem (ENP). Given a set V of n points in the plane, find for each $\tilde{v} \in V$ and $1 \leq l \leq 8$ a nearest neighbor to \tilde{v} in the l th region, if it exists.

We first show that, once the eight-neighbors problem is solved for V , it takes very little extra effort to find an MST on V . To see this, we form E , the set of edges defined by:

$$E = \{\{\tilde{v}, \tilde{u}\} \mid \tilde{v} \in V \text{ and } \tilde{u} \text{ is a nearest neighbor to } \tilde{v} \text{ selected by ENP}\}.$$

We assert that the set of edges E contains an MST on V . As E contains at most $8n$ edges, we can then construct an MST for the sparse graph (V, E) in $O(n \log \log n)$ steps [17], a very small cost.

THEOREM 3.2. *The set of edges E contains an MST on V .*

Proof. Let T be a set of edges that form an MST on V . We will show that, for any edge $\{\tilde{v}, \tilde{w}\}$ that is in T but not in E , we can replace $\{\tilde{v}, \tilde{w}\}$ by an edge in E and still maintain an MST. This would prove the theorem since we can perform this operation on T repeatedly until all edges in T are from E .

Let $\{\tilde{v}, \tilde{w}\}$ be an edge in $T - E$. Assume $\tilde{w} \in R_l(\tilde{v})$. Then $N_l \neq \emptyset$, and there is a nearest neighbor \tilde{u} to \tilde{v} in $N_l(\tilde{v})$ such that $\{\tilde{v}, \tilde{u}\} \in E$. Clearly $\tilde{u} \neq \tilde{w}$ and $d_2(\tilde{v}, \tilde{u}) \leq d_2(\tilde{v}, \tilde{w})$. Let us delete $\{\tilde{v}, \tilde{w}\}$ from T . Then T is separated into two disjoint subtrees with \tilde{v} and \tilde{w} belonging to different components. Now, \tilde{u} and \tilde{w} must be in the same component. For if they were not, $\{\tilde{u}, \tilde{w}\}$ would be a shorter connecting edge for the two subtrees than $\{\tilde{v}, \tilde{w}\}$ by Lemma 3.1, contradicting the fact that T is an MST. Therefore \tilde{u} is in the same subtree as \tilde{w} , and adding the edge $\{\tilde{v}, \tilde{u}\}$ to $T - \{\tilde{v}, \tilde{w}\}$ results in a spanning tree with total weight no greater than that of T . \square

We now proceed to solve the eight-neighbors problem. We will find a nearest neighbor to each point in the first region. The procedure can be simply adapted to find nearest neighbors in the l th region for other l . As demonstrated earlier, the MST problem can be thus solved in a total of $8 \cdot f(n) + O(n \log \log n)$ steps, if the first-region nearest neighbors can be found in $f(n)$ steps.

To study the first regions, it is convenient to tilt the y -axis by 45° clockwise (see Fig. 2). That is, transform the coordinates (x_1, x_2) of a point v into (x'_1, x'_2) , defined by

$$\begin{pmatrix} x'_1 \\ x'_2 \end{pmatrix} = \begin{pmatrix} 1 & -1 \\ 0 & \sqrt{2} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}.$$

In the new coordinates, a point $\tilde{u} = (u'_1, u'_2)$ is in the first region relative to $v = (v'_1, v'_2)$ if and only if $(u'_1 \geq v'_1) \wedge (u'_2 \geq v'_2)$.




FIG. 2. New coordinate system.

For simplicity we assume that all the $2n$ coordinates x'_1, x'_2 of points $\tilde{x} \in V$ are distinct numbers. This restriction shall be removed in the general algorithm in § 3. Let us first sort the points according to their first coordinates x'_1 , and divide them into $s = (n/q)^{1/2}$ consecutive groups each with $\approx qs$ points (Fig. 3), q to be determined later. Then for each of these s groups we sort the points in ascending order of the coordinates x'_2 , and divide them into s consecutive groups with $\approx q$ points each (Fig. 4).

The set V is thus divided into s^2 “cells”. For any $\tilde{v} \in V$, the cells can be classified into three classes by their position relative to \tilde{v} : class 1, cells all of whose points are in $N_1(\tilde{v})$; class 2, cells with no points in $N_1(\tilde{v})$; and class 3, the remaining cells. A useful observation is that the number of cells in class 3 is at most $2 \times s$. This can be understood as follows: If we draw a horizontal and a vertical line through v , only those cells that are “hit” can be in class 3, and there are at most $2 \times s$ of them. We can now try to find a nearest neighbor for \tilde{v} in $N_1(\tilde{v})$ using the following strategy: we examine each cell in turn for cells in class 3, and compute $d_2(\tilde{v}, \tilde{u})$ for all \tilde{u} in the





FIG. 3. Division of points into s groups according to values of x'_1 .


 FIG. 4. Completing the division of V into s^2 cells.

cell; for a cell in class 2, we ignore it; for a cell C in class 1, we compute \tilde{u} and $d_2(\tilde{v}, \tilde{u})$ defined by $d_2(\tilde{v}, \tilde{u}) = \min \{d_2(\tilde{v}, \tilde{x}) | \tilde{x} \in C\}$. A nearest point can now be found by selecting the point \tilde{u} with minimum $d_2(\tilde{v}, \tilde{u})$ from the preceding calculations. The cost is $O(2s \cdot q + \# \text{ of class 1 cells} \times a) = O(2sq + s^2a) = O(n/s + 2an/q)$, where a is the cost of computing $d_2(v_i, C)$ for a cell C of q points. If we have to compute $d_2(\tilde{v}, \tilde{u})$ for each $\tilde{u} \in C$, then $a = O(q)$, and the total cost would be $O(n)$, and we have not made any progress. However, we know from the post-office problem that we can lower a to $\log q$ if we are willing to preprocess the set C (in $O(q^2)$ time). So let us do the following: (i) preprocess every cell C to facilitate the computing of $d_2(\tilde{v}, C)$ (cost $O((n/q) \cdot q^2) = O(nq)$); (ii) for each \tilde{v} , compute the nearest neighbor in the above manner in time $O(n/s + (n/q) \log q)$. The total cost is then $O(nq + n^2/s + (n^2/q) \log q)$. Take $q = n^{1/3}$ and obtain an algorithm that runs in time $O(n^{5/3} \log n)$. This gives an $o(n^2)$ algorithm for finding an MST in two dimensions. We shall generalize the ideas to general k .

4. Reduction of MST to a general GN-problem. We shall prove Theorem 1 for the MST- and GN-problems in this and the next sections. Without loss of generality, we shall assume that the n given points in V are all distinct.

In this section we reduce the finding of MST in E_p^k to a version of the geographic-neighbor problem. We assume that $p \in \{1, 2, \infty\}$ throughout the rest of the paper.

We make E_p^k a vector space by defining $\tilde{x} + \tilde{y} = (x_1 + y_1, x_2 + y_2, \dots, x_k + y_k)$ and $c\tilde{x} = (cx_1, cx_2, \dots, cx_k)$, where c is any real number and x_i, y_i are the components of \tilde{x} and \tilde{y} . We shall refer to any element of E_p^k as a *point* or a *vector*. The j th component of a vector \tilde{z} will be denoted as z_j without further explanation. The *inner product* of two vectors \tilde{x} and \tilde{y} is $\tilde{x} \cdot \tilde{y} = \sum_{i=1}^k x_i y_i$, and the *norm* of \tilde{x} is $\|\tilde{x}\| = (\tilde{x} \cdot \tilde{x})^{1/2}$. A *unit vector* \tilde{x} is a vector with $\|\tilde{x}\| = 1$. Notice that all these definitions are independent of p .

Vectors $\tilde{b}_1, \tilde{b}_2, \dots, \tilde{b}_j$ are *linearly independent* if $\sum_{i=1}^j \lambda_i \tilde{b}_i = 0$ implies all $\lambda_i = 0$. A set of k linearly independent vectors in E_p^k is called a *basis* (of E_p^k). Let $B = \{\tilde{b}_1, \tilde{b}_2, \dots, \tilde{b}_k\}$ be a basis of E_p^k . The *convex cone* of B is $\text{Conv}(B) = \{\sum_{i=1}^k \lambda_i \tilde{b}_i | \lambda_i \geq 0 \text{ for all } i\}$. For any $\tilde{x} \in E_p^k$, the *region* B of \tilde{x} is defined as

$$R(B; \tilde{x}) = \{\tilde{y} | \tilde{y} - \tilde{x} \in \text{Conv}(B)\}.$$

Let V be a set of n distinct vectors in E_p^k . Denote by $N(B, \tilde{v})$ the set $V \cap \{\tilde{u} \mid \tilde{u} \in R(B; \tilde{v}) - \{\tilde{v}\}\}$, for each $\tilde{v} \in V$. We shall say that \tilde{w} is a *geographic neighbor* to \tilde{v} in region B if $\tilde{w} \in N(B; \tilde{v})$ and $d_p(\tilde{w}, \tilde{v}) \leq d_p(\tilde{u}, \tilde{v})$ for all $\tilde{u} \in N(B; \tilde{v})$.

The GGN-problem (general geographic neighbor). Given a basis B and a set V of n distinct vectors in E_p^k , find, for each $\tilde{v} \in V$, a geographic neighbor to \tilde{v} in region B , if one exists.

Notice that this reduces to the GN-problem when $B = \{\tilde{b}_1, \tilde{b}_2, \dots, \tilde{b}_k\}$ with $b_{ij} = \delta_{ij}$. The rest of this section is devoted to showing the following theorem, which states that, if there is a fast algorithm to solve the GGN-problem, then one can solve the MST-problem efficiently.

THEOREM 4.1. *Let $k \geq 2$ be a fixed integer. Suppose there is an algorithm that solves the GGN-problem for n given points in E_p^k in at most $f(n)$ steps. Then a minimum spanning tree for n points in E_p^k can be found in $O(f(n) + n \log \log n)$ steps.*

Define the angle between two nonzero vectors \tilde{x} and \tilde{y} as $\Theta(\tilde{x}, \tilde{y}) = \cos^{-1}(\tilde{x} \cdot \tilde{y} / \|\tilde{x}\| \cdot \|\tilde{y}\|)$, $0 \leq \Theta(\tilde{x}, \tilde{y}) \leq \pi$. For any basis B of E_p^k , the *angular diameter* of B is defined by $\text{Ang}(B) = \sup \{\Theta(\tilde{x}, \tilde{y}) \mid \tilde{x}, \tilde{y} \in \text{Conv}(B)\}$. It can be shown that $\text{Ang}(B) = \max \{\Theta(\tilde{b}_i, \tilde{b}_j) \mid \tilde{b}_i, \tilde{b}_j \in B\}$, although we shall not use that fact.

Let \mathcal{B} be a finite family of basis of E_p^k . We call \mathcal{B} a *frame* if $\bigcup_{B \in \mathcal{B}} \text{Conv}(B) = E_p^k$. The *angular diameter* of a frame \mathcal{B} is given by $\text{Ang}(\mathcal{B}) = \max \{\text{Ang}(B) \mid B \in \mathcal{B}\}$. For example, let $\tilde{b}_1 = (1, 0)$, $\tilde{b}_2 = (-1, 1)$, $\tilde{b}_3 = (0, -1)$, $\tilde{b}_4 = (-\frac{1}{2}, -1)$ as shown in Fig. 5, then




FIG. 5. Illustration of “basis” and “frame”.

$B_1 = \{\tilde{b}_1, \tilde{b}_2\}$, $B_2 = \{\tilde{b}_2, \tilde{b}_3\}$, $B_3 = \{\tilde{b}_4, \tilde{b}_1\}$ are bases of E_p^2 , and $\mathcal{B} = \{B_1, B_2, B_3\}$ a frame; $\Theta(B_1) = \Theta(B_2) = 3\pi/4$, $\Theta(B_3) = 2\pi/3$, and $\Theta(\mathcal{B}) = 3\pi/4$.

Intuitively, the convex cone of a basis B has a “narrow” angular coverage if $\text{Ang}(B)$ is small. The following result asserts that a frame exists in which every basis is narrow, and such a frame can be constructed.

LEMMA 4.2. *For any $0 < \psi < \pi$, one can construct in finite steps a frame \mathcal{B} of E_p^k such that $\text{Ang}(\mathcal{B}) < \psi$.*

Proof. See Appendix. \square

We consider the following MST-algorithm. Let us construct a frame \mathcal{B} of E_p^k such that $\text{Ang}(B) < \sin^{-1}(\frac{1}{2}k^{-(1/2+1/p)})$. Next, for each $B \in \mathcal{B}$, we solve the GGN-problem—for each $\tilde{v} \in V$, find a geographic neighbor \tilde{u} to \tilde{v} in region B if it exists—and form the set $E(B)$, the collection of all such edges $\{\tilde{u}, \tilde{v}\}$. Clearly, $|\bigcup_{B \in \mathcal{B}} E(B)| \leq n \cdot |\mathcal{B}| = O(n)$. We now claim that $\bigcup_{B \in \mathcal{B}} E(B)$ contains an MST on V . If this is true, then we can find an MST in an additional $O(n \log \log n)$ steps. The total time taken by the MST algorithm is then $O(f(n) + n \log \log n)$. It remains to prove the following result.

LEMMA 4.3. $\bigcup_{B \in \mathcal{B}} E(B)$ contains an MST on V .

Proof. The proof is almost identical to the proof of Theorem 3.2, except that we need to establish the next lemma. \square .

LEMMA 4.4. Let $\tilde{x}, \tilde{y}, \tilde{z}$ in E_p^k satisfy $\Theta(\tilde{x} - \tilde{z}, \tilde{y} - \tilde{z}) < \sin^{-1}(\frac{1}{2}k^{-(1/2+1/p)})$, then $d_p(\tilde{x}, \tilde{y}) < \max\{d_p(\tilde{y}, \tilde{z}), d_p(\tilde{x}, \tilde{z})\}$.

Proof. Use α, β, γ to denote angles as shown in Fig. 6. By assumption,

$$(1) \quad \sin \alpha < \frac{1}{2}k^{-(1/2+1/p)}.$$




FIG. 6. Illustration for the proof of Lemma 4.4.

Without loss of generality, assume that $\alpha + \beta > \pi/2$. Let \tilde{w} be the projection of \tilde{y} on the segment from \tilde{z} to \tilde{x} . By the triangle inequality satisfied by metric d_p , we have

$$\begin{aligned} d_p(\tilde{z}, \tilde{w}) + d_p(\tilde{w}, \tilde{y}) &\geq d_p(\tilde{y}, \tilde{z}), \\ d_p(\tilde{x}, \tilde{w}) + d_p(\tilde{w}, \tilde{y}) &\geq d_p(\tilde{x}, \tilde{y}). \end{aligned}$$

Thus,

$$(2) \quad d_p(\tilde{z}, \tilde{w}) + d_p(\tilde{x}, \tilde{w}) \geq d_p(\tilde{x}, \tilde{y}) + (d_p(\tilde{y}, \tilde{z}) - 2d_p(\tilde{w}, \tilde{y})).$$

But since \tilde{w} is on the segment \tilde{z} to \tilde{x} , we have $d_p(\tilde{x}, \tilde{z}) = d_p(\tilde{z}, \tilde{w}) + d_p(\tilde{x}, \tilde{w})$. Therefore, if we can further show that

$$(3) \quad d_p(\tilde{y}, \tilde{z}) - 2d_p(\tilde{w}, \tilde{y}) > 0,$$

then (2) implies $d_p(\tilde{x}, \tilde{z}) > d_p(\tilde{x}, \tilde{y})$, proving the lemma.

To prove formula (3), we notice that for any positive l , and \tilde{u}, \tilde{v} in E_l^k ,

$$(4) \quad k^{1/l} \max_i |\tilde{u}_i - \tilde{v}_i| \geq d_l(\tilde{u}, \tilde{v}) \geq \max_i |\tilde{u}_i - \tilde{v}_i|.$$

This leads to

$$(5) \quad k^{1/p} \|\tilde{u} - \tilde{v}\| \geq d_p(\tilde{u}, \tilde{v}) \geq k^{-1/2} \|\tilde{u} - \tilde{v}\|.$$

In particular,

$$(6) \quad \begin{aligned} d_p(\tilde{y}, \tilde{w}) &\leq k^{1/p} \|\tilde{y} - \tilde{w}\|, \\ d_p(\tilde{y}, \tilde{z}) &\geq k^{-1/2} \|\tilde{y} - \tilde{z}\|. \end{aligned}$$

Now, clearly by (1),

$$(7) \quad \|\tilde{y} - \tilde{w}\| = (\sin \alpha) \|\tilde{y} - \tilde{z}\| < \frac{1}{2}k^{-(1/2+1/p)} \|\tilde{y} - \tilde{z}\|.$$

Formula (3) follows from (6) and (7). \square

5. An algorithm for the general geographic neighbor problem.

5.1. An outline. As shown in the preceding section, the MST-problem can be reduced to the GGN-problem, and the GN-problem is a special case of the GGN-

problem. In this section, we shall give an asymptotically fast algorithm for the GGN-problem, which completes the proof of Theorem 1.

Given a basis B and a set V of n points in E_p^k , the algorithm works in two phases.

Preprocessing phase. (A) Partition V in $O(kn \log n)$ steps into $r = \lceil n/q \rceil$ subsets V_1, V_2, \dots, V_r , each with at most q points (q to be determined later). The division will be such that, for any $\tilde{x} \in E_p^k$, all but a fraction $r^{-1/k}$ of the subsets V_j have the property that the entire set V_j is either in region B of \tilde{x} or outside of region B .

(B) Preprocess each V_j in $O(q^{b(k)})$ steps such that, for any new point $\tilde{x} \in E^k$, a nearest point \tilde{u} in V_j can be found in $O(\log q)$ steps.

Finishing phase. (C) For each $\tilde{v} \in V$, we find a geographic neighbor in region B as follows. We examine the r sets V_1, V_2, \dots, V_r in turn. For each V_j , we perform a test which puts V_j into one of the three categories. A category-1 V_j has all its points in region B of \tilde{v} , a category-2 V_j has all its points outside of region B . The nature of a category-3 V_j is unimportant, except that there are at most r^{1-k-1} V_j in this category; we consider the V_j that contains \tilde{v} itself to be of category 3 independent of the above division. As we shall see later, the test will be easy to carry out, in fact in $O(k)$ time per test. For a category-1 V_j , we find a nearest \tilde{w} in V_j in $O(\log q)$ time. For a category-2 V_j , nothing need be done. For a category-3 V_j , we find a nearest $\tilde{w} (\neq \tilde{v}) \in V_j$ in region B , if it exists, by finding all the $\tilde{z} \in V_j$ that are in region B and computing and comparing $d_p(\tilde{z}, \tilde{v})$ for all such \tilde{z} . Call \tilde{w} a candidate from V_j . After all the V_j have been so processed, we compare $d_p(\tilde{w}, \tilde{v})$ for all the candidates \tilde{w} obtained (at most r of them), and find a nearest one \tilde{u} to \tilde{v} . This \tilde{u} is the geographic neighbor we seek for \tilde{v} . Return “nonexistent” if no candidate \tilde{w} exists from any V_j .

In the above description, three points need further elaboration: how step (A) is accomplished, how we check a subset V_i for its category, and how q is chosen. We shall deal with the first two points in § 5.2, and the last point in § 5.3.

5.2. A set partition theorem. We shall show that step (A) of the preprocessing phase in § 5.1 can be accomplished. The key is the following result in Yao and Yao [20]. For completeness, a proof is included.

For any finite set F of points in E^k , let $\text{high}_l(F) = \max \{x_l \mid \tilde{x} \in F\}$ and $\text{low}_l(F) = \min \{x_l \mid \tilde{x} \in F\}$, for $1 \leq l \leq k$.

LEMMA 5.1 [20].² Let q and k be positive integers, and F a set of n points in E^k . Then, in $O(kn \log n)$ steps, the following can be done.

- (i) F is partitioned into $r = \lceil n/q \rceil$ sets F_1, F_2, \dots, F_r , each with at most q points,
- (ii) The $2kr$ numbers $\text{high}_l(F_i), \text{low}_l(F_i)$, $1 \leq i \leq r$, and $1 \leq l \leq k$, are computed,
- (iii) The partition satisfies the condition that, for any $\tilde{y} \in E^k$, there exist at most $k \lceil r^{1/k} \rceil^{k-1}$ sets F_i such that there exists an l with $\text{low}_l(F_i) < y_l \leq \text{high}_l(F_i)$.

Proof. We shall prove Lemma 5.1 for the case $k = 3$; the extension to general k is obvious. For the moment, let us assume further that $n = qm^3$ for some integer m . We use the following procedure to partition F .

(a) Sort the points of F in ascending order according to the first components into a sequence $\tilde{x}_1, \tilde{x}_2, \dots, \tilde{x}_n$. Divide the sorted sequence into m consecutive parts of equal size, that is,

$$G_1 = \{\tilde{x}_j \mid 1 \leq j \leq n/m\}, \quad G_2 = \{\tilde{x}_j \mid n/m + 1 \leq j \leq 2n/m\}, \quad \dots, \quad G_m.$$

² This lemma was proved in [20] with $q = n^{1/k}$.

(b) For each $1 \leq i \leq m$, sort the points in G_i according to the 2nd components; divide the sorted sequence of G_i into m consecutive parts of equal size, $G_{i1}, G_{i2}, \dots, G_{im}$.

(c) For each $1 \leq i, j \leq m$, sort the points in G_{ij} according to their 3rd components; divide the sorted sequence of G_{ij} into m consecutive parts of equal size, $G_{ij1}, G_{ij2}, \dots, G_{ijm}$.

(d) Rename the m^3 sets G_{ijt} as F_1, F_2, \dots, F_r , where $r = n/q = m^3$.

(e) Compute $\text{high}_l(F_i), \text{low}_l(F_i)$ for $1 \leq i \leq r, 1 \leq l \leq 3$ according to the definitions. The above procedure takes $O(n \log n)$ steps, and each F_i contains exactly q points. It remains to show that property (iii) in Lemma 5.1 is satisfied.

Let $\tilde{y} \in E^3$. We shall prove that, for each $1 \leq l \leq 3$, there are at most m^2 F_i with $\text{low}_l(F_i) < y_l \leq \text{high}_l(F_i)$. The proof is based on the following properties of the partition, where $1 \leq i, j \leq m$:

$$(5.1) \quad \text{low}_1(G_1) \leq \text{high}_1(G_1) \leq \text{low}_1(G_2) \leq \text{high}_1(G_2) \\ \leq \dots \leq \text{low}_1(G_m) \leq \text{high}_1(G_m),$$

$$(5.2) \quad \text{low}_2(G_{i1}) \leq \text{high}_2(G_{i1}) \leq \text{low}_2(G_{i2}) \leq \text{high}_2(G_{i2}) \\ \leq \dots \leq \text{low}_2(G_{im}) \leq \text{high}_2(G_{im}),$$

$$(5.3) \quad \text{low}_3(G_{ij1}) \leq \text{high}_3(G_{ij1}) \leq \text{low}_3(G_{ij2}) \leq \text{high}_3(G_{ij2}) \\ \leq \dots \leq \text{low}_3(G_{ijm}) \leq \text{high}_3(G_{ijm}).$$

For $l = 1$, according to (5.1), there is at most one j such that

$$\text{low}_1(G_j) < y_1 \leq \text{high}_1(G_j).$$

Thus, only the m^2 G_{jts} , $1 \leq t, s \leq m$, can have $\text{low}_1(G_{jts}) < y_1 \leq \text{high}_1(G_{jts})$. This proves our assertion for $l = 1$. We now prove the case for $l = 2$. For each i , by (5.2), there is at most one j such that $\text{low}_2(G_{ij}) < y_2 \leq \text{high}_2(G_{ij})$. Thus, for each i , only the m G_{ijt} , $1 \leq t \leq m$, may have $\text{low}_2(G_{ijt}) < y_2 \leq \text{high}_2(G_{ijt})$. Therefore, at most m^2 G_{ijt} can have $\text{low}_2(G_{ijt}) < y_2 \leq \text{high}_2(G_{ijt})$. A similar proof works for $l = 3$, making use of formula (5.3).

This proves that, when $k = 3$, and $n = qr = qm^3$ for some integer m , Lemma 5.1 is true. We now drop the restriction on n (still $k = 3$). In this situation, $r = \lceil n/q \rceil$. Let $m = \lceil r^{1/3} \rceil$, and use the same procedure. At most $3m^2$ G_{ijt} will satisfy (iii) by the same proof. This completes the proof for $k = 3$. \square

We now extend the above result. Let $B = \{\tilde{b}_1, \tilde{b}_2, \dots, \tilde{b}_k\}$ be a basis of E^k ; for any $\tilde{x} \in E^k$, we shall define a k -tuple $(x'_1, x'_2, \dots, x'_k)$ by $\tilde{x} = \sum_{i=1}^k x'_i \tilde{b}_i$. For any finite set F of points, define for each $1 \leq l \leq k$,

$$\text{high}_l(B; F) = \max \{x'_l \mid \tilde{x} \in F\}, \\ \text{low}_l(B; F) = \min \{x'_l \mid \tilde{x} \in F\}.$$

THEOREM 5.2. *Let q, n, k ($q, k \leq n$) be positive integers, B a basis of E^k , and V a set of n points in E^k . Then, in $O(kn \log n + k^2 n + k^3)$ steps, we can accomplish the following:*

- (i) V is partitioned into $r = \lceil n/q \rceil$ sets V_1, V_2, \dots, V_r , each with at most q points.
- (ii) The $2kr$ numbers $\text{high}_l(B, V_i), \text{low}_l(B, V_i)$, $1 \leq i \leq r, 1 \leq l \leq k$, are computed.

Furthermore, the partition satisfies the condition:

(iii) For any k -tuple of numbers (y_1, y_2, \dots, y_k) , there exist at most $k \lceil r^{1/k} \rceil^{k-1} V_i$ such that there exists an l that satisfies:

$$\text{low}_l(B; V_i) < y_l \leq \text{high}_l(B; V_i).$$

Before proving this theorem, let us check that this partition fulfills the requirements of step (i) in the preprocessing phase (see § 5.1).

LEMMA 5.3. A point \tilde{y} is in the region B to \tilde{x} , i.e., $\tilde{y} \in R(B; \tilde{x})$, if and only if $y'_l \geq x'_l$ for all $1 \leq l \leq k$.

Proof. The lemma follows from the equation $\tilde{y} - \tilde{x} = \sum_{l=1}^k (y'_l - x'_l) \tilde{b}_l$. \square

LEMMA 5.4. If $\tilde{x} \in E^k$, B a basis, and F a finite set of points in E^k , then either

(i) $x'_l \leq \text{low}_l(B; F)$ for all $1 \leq l \leq k$, in which case all points in F are in region B to \tilde{x} , or

(ii) $\exists l, x'_l > \text{high}_l(B; F)$, in which case none of the points in F is in region B to \tilde{x} , or

(iii) none of the above; there exists an l such that $\text{low}_l(A; F) < x'_l \leq \text{high}_l(B; F)$.

Proof. This is an immediate consequence of Lemma 5.3. \square

There are two consequences of Lemma 5.4 of interest to us. Firstly, the lemma shows that the requirements of step (A) in § 5.1 are satisfied. For any \tilde{x} , a V_j such that neither all points of V_j are in $R(B; \tilde{x})$ nor none are in $R(B; \tilde{x})$ must satisfy the condition that $\text{low}_l(B, V_j) < x'_l \leq \text{high}_l(B; V_j)$ for some l , because of Lemma 5.4. By Theorem 5.2, there are at most about $r^{1-1/k}$ such V_j . This proves the claim. Secondly, Lemma 5.4 gives a simple way to detect most of the V_j that satisfy $V_j \subseteq R(B; \tilde{x})$ or $V_j \cap R(B; \tilde{x}) = \emptyset$. Namely, compare x'_l with $\text{high}_l(B; V_j)$ and $\text{low}_l(B; V_j)$ for all l , and determine whether case (i), (ii), or (iii) applies in Lemma 5.4. The test only takes $O(k)$ for each i and j , can be conveniently used in step (C) in the procedure in § 5.1.

We now turn to the proof of Theorem 5.2.

Proof of Theorem 5.2. Let M be the $k \times k$ matrix (b_{ij}) (recall that $\tilde{b}_i = (b_{i1}, b_{i2}, \dots, b_{ik})$), and M^{-1} be its inverse. We use the following procedure to partition V .

- (1) Compute M^{-1} in $O(k^3)$ steps (see, e.g., [1]).
- (2) Compute, for each $\tilde{x} \in V$, the k -tuple $(x'_1, x'_2, \dots, x'_k)$ by $(x'_1, x'_2, \dots, x'_k) = (x_1, x_2, \dots, x_k) \cdot M^{-1}$. This takes $O(k^2 n)$ steps.
- (3) Consider the set $F = \{(x'_1, x'_2, \dots, x'_k) | \tilde{x} \in V\}$. We now use the procedure in Lemma 5.1 to divide F into r parts F_1, F_2, \dots, F_r . Let V_i be the subset of V obtained from F_i by replacing every (x'_1, \dots, x'_k) by the corresponding \tilde{x} .
- (4) Set $\text{high}_l(B; V_i) \leftarrow \text{high}_l(F_i)$, and $\text{low}_l(B; V_i) \leftarrow \text{low}_l(F_i)$.

The procedure clearly takes $O(kn \log n + k^2 n + k^3)$ steps. The quantities $\text{high}_l(B; V_i)$ and $\text{low}_l(B; V_i)$ are correctly computed by their definitions. Items (i) and (ii) in Theorem 5.2 are obviously true, and (iii) is true because of the properties of $\text{high}_l(F_i)$, $\text{low}_l(F_i)$ stated in Lemma 5.1. \square

5.3. Finishing the proof. We now analyze the running time of the algorithm for fixed k and choose q . The preprocessing phase takes time $O(n \log n + r \cdot q^{b(k)})$. In the finishing phase, the running time is dominated by the search for candidates \tilde{w} , which is of order $n[(\# \text{ of category-1 } V_j) \cdot \log q + (\# \text{ of category-3 } V_j) \cdot q]$. The last expression is bounded by $n(r \log q + r^{1-k-1} \cdot q)$. The total running time of the algorithm is thus $O(n \log n + r \cdot q^{b(k)} + nr \log q + nqr^{1-k-1})$. Remembering that $b(k) = 2^{k+1}$ and $r = O(n/q)$, we optimize the expression by choosing $q \approx (n \log n)^{a(k)}$. This gives a time $O(n^{2-a(k)} (\log n)^{1-a(k)})$. The improved time bound for the special case $k = 3, p = 2$ can be similarly obtained.

6. Discussions. We have shown that, for fixed k and $p \in \{1, 2, \infty\}$, there are $o(n^2)$ -time algorithms for a number of geometric problems in E_p^k , including the minimum spanning tree problem. We shall now argue that, when $p \in \{2, \infty\}$, $o(kn^2)$ algorithms exist for all k and n . As are typical for results under fixed k assumptions, the algorithms given in the paper have $o(n^2)$ time bounds when k is allowed to grow slowly with n . In fact, a close examination shows that, if $k \leq \frac{1}{2} \log \log n$, the algorithms still run in time $o(n^2)$. For $k > \frac{1}{2} \log \log n$, it can be shown [18] that the computation of the distances between all points can be done in $o(kn^2)$ time when $p \in \{2, \infty\}$. Since all problems considered in this paper have $O(n^2)$ algorithms once all the distances are known, the previous statement provides algorithms that run in time $o(kn^2)$.

The efficiency of our algorithms is dependent on the solution to the post-office problem (or its farthest-point analogue). For example, suppose the nearest-point query could be answered in $O(\log n)$ time after an $O(n^\beta)$ -time preprocessing, $\beta \geq 2$. A simple adaptation of the algorithm would give an $O(n^{2-\beta-1}(\log n)^{1-\beta-1})$ -time solution to the NFN-problem, which in turn implies an $O((n \log n)^{2-\beta-1})$ -time solution to the MST-problem (see the remark at the end of § 2). If $1 < \beta < 2$, the following modification would also give an $O(n^{2-\beta-1}(\log n)^{1-\beta-1})$ -algorithm for the NFN-problem (and hence an $O((n \log n)^{2-\beta-1})$ -algorithm for finding MST). We first divide V into $r \approx n/(n \log n)^{\beta-1}$ blocks B_1, B_2, \dots as before. Each block is preprocessed, and for each \tilde{x} , a nearest point in every block not containing x is found. Now, for every point $\tilde{x} \in B_i$, we need to find for it a nearest “foreign” neighbor in B_i . Instead of using brute force (computing the distance from each $\tilde{x} \in B_i$ to every other point in B_i) as was done previously, we divide B_i into r subblocks, preprocess each subblock, and find for \tilde{x} a nearest point in every subblock in B_i . To compute a nearest foreign neighbor to \tilde{x} in the subblock containing \tilde{x} , we shall again break the subblocks. This process continues until the size of the subblocks is less than n^δ , where $\delta = 1 - \beta^{-1}$, at which point we compute all distances between points in the same subblock. During the above process, we have located, for each \tilde{x} , a set of points containing a nearest foreign neighbor \tilde{u} to \tilde{x} . It is then simple to locate such a \tilde{u} . This is a brief outline of an $O(n^{2-\beta-1}(\log n)^{1-\beta-1})$ -algorithm for NFN-problems, $1 < \beta < 2$. However, it seems unlikely that a nearest-point query can be answered in $O(\log n)$ time with an $O(n^\beta)$ -preprocessing, $\beta < 2$, when $k \geq 3$.

We conclude this paper with the following open problems.

- (1) Improve the bounds obtained in this paper.
- (2) Analyze the performance of new or existing fast heuristic algorithms for MST-problems. For example, can one show that the AMST-algorithm in [2] always constructs a spanning tree with length at most 5% over the true MST?
- (3) Prove bounds on average running time of MST-algorithms for some natural distributions.
- (4) Extend results in this paper to L_p -metric for general p .

Appendix. The existence and construction of “narrow” frames—Proof of Lemma 4.2.

LEMMA 4.2. *For any $0 < \psi < \pi$, one can construct in finite steps a frame \mathcal{B} of E_p^k such that $\text{Ang}(\mathcal{B}) < \psi$.*

As the discussion is independent of p , we shall use E^k instead of E_p^k .

We begin with the concept of a “simplex” familiar in Topology (see, e.g., [10]). Let $\tilde{p}_0, \tilde{p}_1, \dots, \tilde{p}_j$ be $j+1$, $0 \leq j \leq k$, points in E^k , where the vectors $\tilde{p}_i - \tilde{p}_0$, $1 \leq i \leq j$, are linearly independent. We shall call the set $\{\sum_{i=0}^j \lambda_i \tilde{p}_i \mid \lambda_i \geq 0 \text{ for all } i, \text{ and } \sum_i \lambda_i = 1\}$ a (geometric) j -simplex in E^k , denoted by $\langle \tilde{p}_0, \tilde{p}_1, \dots, \tilde{p}_j \rangle$. Informally, it is the convex

hull formed by vertices $\tilde{p}_0, \tilde{p}_1, \dots, \tilde{p}_i$ on the minimal linear subspace containing them (see Fig. A). The *diameter* of a simplex s is $\text{diam}(s) = \sup \{\|\tilde{x} - \tilde{y}\| \mid \tilde{x}, \tilde{y} \in s\}$.

The following two lemmas give the connection between simplices and bases. Let $\hat{\varepsilon}$ be a k -tuple $(\varepsilon_1, \varepsilon_2, \dots, \varepsilon_k)$, where $\varepsilon_i \in \{-1, 1\}$ for all i . Denote by $H(\hat{\varepsilon})$ the hyperplane $\{x \mid \sum_i \varepsilon_i x_i = 1\}$ in E^k .




FIG A. A 2-simplex in E^2

LEMMA A1. Let $s = \langle \hat{p}_0, \hat{p}_1, \dots, \hat{p}_{k-1} \rangle$ be a $(k-1)$ -simplex in E^k , where $\tilde{p}_i \in H(\hat{\varepsilon})$ for every i . Then the set $B(s) = \{\tilde{p}_0, \tilde{p}_1, \dots, \tilde{p}_k\}$ is a basis. Furthermore, the angle $\varphi = \text{Ang}(B(s))$ satisfies $\cos \varphi \geq 1 - \frac{1}{2}k(\text{diam}(s))^2$.

Proof. Suppose $\sum_{i=0}^{k-1} \lambda_i \tilde{p}_i = 0$. We shall show that $\lambda_i = 0$ for all i . If $\sum_{i=0}^{k-1} \lambda_i = 0$, then $\sum_{i=1}^{k-1} \lambda_i (\tilde{p}_i - \tilde{p}_0) = \sum_{i=0}^{k-1} \lambda_i \tilde{p}_i = 0$. This implies $\lambda_i = 0$ for all i , by the definition of simplex. If $\sum_{i=0}^{k-1} \lambda_i = \Lambda \neq 0$, then $\tilde{v} = \sum_{i=0}^{k-1} (\lambda_i / \Lambda) \tilde{p}_i = 0$. But it is easy to check that $v \in H(\hat{\varepsilon})$, a contradiction.

We have thus shown that $B(s)$ is a basis. To prove the rest of the lemma, let \tilde{x} and \tilde{y} be any two nonzero vectors in $\text{Conv}(B(s))$. We shall prove that $\cos \theta(\tilde{x}, \tilde{y}) \geq 1 - \frac{1}{2}k(\text{diam}(s))^2$. Without loss of generality, we can assume that $\tilde{x}, \tilde{y} \in s$. Then

$$\begin{aligned} (\text{diam}(s))^2 &\geq (\tilde{x} - \tilde{y}) \cdot (\tilde{x} - \tilde{y}) = \|\tilde{x}\|^2 + \|\tilde{y}\|^2 - 2\|\tilde{x}\| \cdot \|\tilde{y}\| \cos \theta(\tilde{x}, \tilde{y}) \\ &\geq 2\|\tilde{x}\| \cdot \|\tilde{y}\| (1 - \cos \theta(\tilde{x}, \tilde{y})). \end{aligned}$$

It follows that

$$(A1) \quad \cos \theta(\tilde{x}, \tilde{y}) \geq 1 - \frac{(\text{diam}(s))^2}{2\|\tilde{x}\| \cdot \|\tilde{y}\|}.$$

As can be easily verified, $\tilde{x}, \tilde{y} \in H(\hat{\varepsilon})$, which implies

$$\|\tilde{x}\|^2 = \sum_i x_i^2 \geq \frac{1}{k} \left(\sum_i \varepsilon_i x_i \right)^2 = \frac{1}{k}.$$

Therefore, $\|\tilde{x}\| \geq 1/\sqrt{k}$, and similarly $\|\tilde{y}\| \geq 1/\sqrt{k}$. Formula (A1) then implies

$$\cos \theta(\tilde{x}, \tilde{y}) \geq 1 - \frac{k}{2} (\text{diam}(s))^2.$$

This proves Lemma A1. \square

We shall use $B(s)$ to denote the basis corresponding to simplex s .

LEMMA A2. Let $s \subseteq H(\hat{\varepsilon})$ be a simplex, \mathcal{S} a finite collection of simplices, and $s = \bigcup_{s' \in \mathcal{S}} s'$. Then $\text{Conv}(B(s)) = \bigcup_{s' \in \mathcal{S}} \text{Conv}(B(s'))$.

Proof. It is easy to see that $\text{Conv}(B(s)) \supseteq \bigcup_{s' \in \mathcal{S}} \text{Conv}(B(s'))$. To prove the converse, let $s = \langle \tilde{p}_0, \tilde{p}_1, \dots, \tilde{p}_{k-1} \rangle$, where each $\tilde{p}_i \in H(\hat{\varepsilon})$. If a point $\tilde{u} \in \text{Conv}(B(s))$, then $\tilde{u} = \sum_{i=0}^{k-1} \lambda_i \tilde{p}_i$, where $\lambda_i \geq 0$. We shall prove that $\tilde{u} \in \text{Conv}(B(s'))$ for some $s' \in \mathcal{S}$. It is trivial if $\tilde{u} = 0$. Otherwise, the point $(1/\sum_i \lambda_i) \tilde{u} \in s = \bigcup_{s' \in \mathcal{S}} s'$, and hence $(1/\sum_i \lambda_i) \tilde{u} \in s'$ for some $s' \in \mathcal{S}$. This implies $\tilde{u} \in \text{Conv}(B(s'))$. \square

The above lemmas suggest that we may try to construct a frame with narrow bases, by first constructing a family of simplices all with small diameters. We use the following scheme:

Let \tilde{e}_i denote the unit vector in E^k , whose i th component is 1 and all others are 0.

For each of the 2^k k -tuples $\hat{e} = (\varepsilon_1, \varepsilon_2, \dots, \varepsilon_k)$, where $\varepsilon_i = \pm 1$, do the following.

(a) Let $s = \langle \varepsilon_1 \tilde{e}_1, \varepsilon_2 \tilde{e}_2, \dots, \varepsilon_k \tilde{e}_k \rangle$. (Clearly $s \subseteq H(\hat{e})$.)

(b) Construct a finite family \mathcal{S} of simplices all contained in $H(\hat{e})$ such that $s = \bigcup_{s' \in \mathcal{S}} s'$ and $\text{diam}(s') < (2(1 - \cos \psi)/k)^{1/2}$ for all $s' \in \mathcal{S}$.

(c) Form $B(s')$ for all $s' \in \mathcal{S}$.

The collection \mathcal{B} of all the $B(s')$ constructed this way is clearly a frame because of Lemma A2. Using Lemma A1, it is easy to verify that $\text{Ang}(B(s')) < \psi$ for all s' . Thus, such a construction would give a frame satisfying the conditions in Lemma 4.2. It remains to show that step (b) above can be carried out.

A procedure in topology ([10, p.209, Thm. 5–20]), known as *barycentric subdivision*, guarantees that step (b) can be accomplished in a finite number of steps. For completeness, we shall give a brief description below.

There is a basis procedure, called *first barycentric subdivision* (FBS), which, for a given j -simplex s , constructs in finite steps a family \mathcal{S} of simplices such that $s = \bigcup_{s' \in \mathcal{S}} s'$ and $\max_{s' \in \mathcal{S}} (\text{diam}(s')) \leq (j/(j+1))(\text{diam}(s))$. If we apply this procedure iteratively, at each iteration we apply FBS to every simplex present, then all the simplices will have a diameter less than any prescribed positive number after enough iterations. This then constitutes a procedure for step (b).

Finally, we describe the FBS procedure. For a proof that it produces simplices with the desired properties, see [10]. Let $s = \langle \tilde{p}_0, \tilde{p}_1, \dots, \tilde{p}_j \rangle$, the point $\tilde{c}(s) = (1/(j+1)) \sum_{i=0}^j \tilde{p}_i$ is called the *centroid* of simplex s . For any t distinct integers $0 \leq i_1, i_2, \dots, i_t \leq j$, let $\tilde{p}_{i_1 i_2 \dots i_t} = \tilde{c}(\langle \tilde{p}_{i_1}, \tilde{p}_{i_2}, \dots, \tilde{p}_{i_t} \rangle)$. For each $\sigma = (i_0, i_1, \dots, i_j) \in \Sigma$, where Σ is the set of all permutations of $(0, 1, 2, \dots, j)$, let $s'(\sigma)$ denote the simplex $\langle \tilde{p}_{\sigma_0}, \tilde{p}_{\sigma_1}, \dots, \tilde{p}_{\sigma_j} \rangle$ with $\sigma_t = i_0, i_1, \dots, i_t$. The FBS of s is defined by

$$\mathcal{S} = \{s'(\sigma) \mid \sigma \in \Sigma\}.$$

REFERENCES

- [1] A. V. AHO, J. E. HOPCROFT AND J. D. ULLMAN, *The Design and Analysis of Computer Algorithms*, Addison-Wesley, Reading, MA, 1974.
- [2] J. L. BENTLEY AND J. H. FRIEDMAN, *Fast algorithms for constructing minimal spanning trees in coordinate spaces*, Rep. STAN-CS-75-529, Stanford Computer Science Department, Stanford, CA, January 1976.
- [3] J. L. BENTLEY AND M. I. SHAMOS, *Divide-and-conquer in multidimensional space*, in Proc. 8th Annual ACM Symposium on Theory of Computing, 1976, pp. 220–230.
- [4] C. BERGE AND A. GHOUILA-HOURI, *Programming, Games, and Transportation Networks*, John Wiley, New York, 1965.
- [5] R. C. BUCK, *Partition of space*, Amer. Math. Monthly, 50 (1943), pp. 541–544.
- [6] D. CHERITON AND R. E. TARJAN, *Finding minimum spanning trees*, this Journal, 5 (1976), pp. 724–742.
- [7] E. W. DIJKSTRA, *A note on two problems in connection with graphs*, Numer. Math., 1 (1959), pp. 269–271.
- [8] D. DOBKIN AND R. J. LIPTON, *Multidimensional search problems*, this Journal, 5 (1976), pp. 181–186.
- [9] R. O. DUDE AND P. E. HART, *Pattern Classification and Science Analysis*, John Wiley, New York, 1973.
- [10] J. G. HOCKING AND G. S. YOUNG, *Topology*, Addison-Wesley, Reading, MA, 1961.
- [11] A. KERSCHENBAUM AND R. VAN SLYKE, *Computing minimum spanning trees efficiently*, in Proc. 25th Annual Conference of the ACM, 1972, pp. 518–527.

- [12] D. E. KNUTH, *The Art of Computer Programming, Vol. 3, Sorting and Searching*, Addison-Wesley, Reading, MA, 1973.
- [13] J. B. KRUSKAL, JR., *On the shortest spanning subtree of a graph and the travelling salesman problem*, Problem. Amer. Math. Soc., 7 (1956), pp. 48–50.
- [14] R. C. PRIM, *Shortest connection networks and some generalizations*, Bell System Tech. J., 36 (1957), pp. 1389–1401.
- [15] M. I. SHAMOS, *Geometric complexity*, Proc. 7th Annual ACM Symposium on Theory of Computing, 1975, pp. 224–233.
- [16] M. I. SHAMOS AND D. J. HOEY, *Closest-point problems*, Proc. 16th Annual IEEE Symposium on Foundations of Computer Science, 1975, pp. 151–162.
- [17] A. C. YAO, *An $O(|E| \log \log |V|)$ algorithm for finding minimum spanning trees*, Inform. Process. Lett., 4 (1975), pp. 21–23.
- [18] ———, *On computing the distance matrix of n points in k dimensions*, IBM San Jose Research Center Technical Report, to appear.
- [19] ———, *On the preprocessing cost in multidimensional search*, IBM San Jose Research Center Technical Report, to appear.
- [20] A. C. YAO AND F. F. YAO, *On computing the rank function for a set of vectors* Rep. UIUCDCS-R-75-699, Computer Science Department, University of Illinois, Illinois, February, 1975.
- [21] C. T. ZAHN, *Graph-theoretical method for detecting and describing gestalt clusters*, IEEE Trans. Comput., C-20 (1970), pp. 68–86.